

## EKS Launch Template Configuration Cheat Sheet

Scenario	update_default_version	version	Behavior	When to Use
Initial Creation	0	"1" or "\$Default"	Creates LT v1, nodes use v1	First-time setup
Safe Updates	0	aws_launch_template.example.latest_version	Creates new LT version, rolls out nodes	✅ Recommended for all updates
Dangerous	1	"\$Latest"	❌ Race condition, suspended nodes	❌ Never use this
No Rollout	0	"5" (pinned)	Creates new LT version but no rollout	Testing changes safely
Broken	1	"5" (pinned)	Wasted change: updates default but no rollout	❌ Never use this

### When we have new changes

- Do this:
  - update\_default\_version = false
  - version = aws\_launch\_template.example.latest\_version
- Not This:
  - update\_default\_version = true
  - version = \$Latest

### Rollout Behavior:

Configuration	Creates New Version	Rolls Out Nodes	Safe
update_default_version=false + version=latest_version	✅ Yes	✅ Yes	✅ Yes
update_default_version=false + version="5"	✅ Yes	❌ No	✅ Yes
update_default_version=true + version="5"	✅ Yes	❌ No	❌ No
update_default_version=true + version="\$Latest"	✅ Yes	✅ Yes	❌ No

### Recovery from suspended state

# 1. Check status

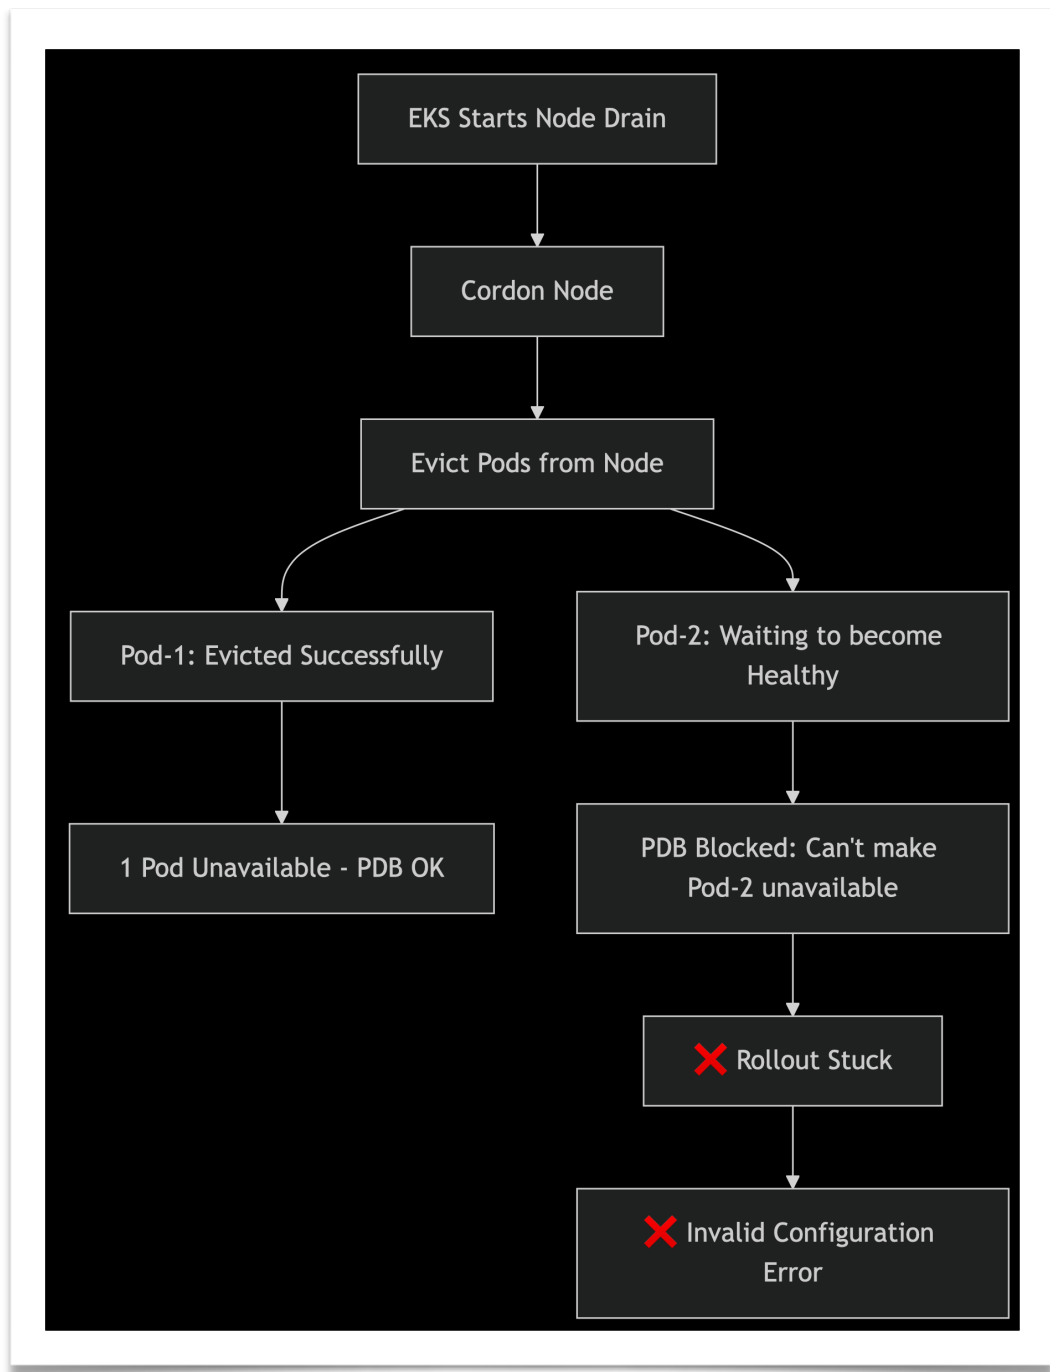
```
aws eks describe-nodegroup --cluster-name my-cluster --nodegroup-name my-ng
```

# 2. If suspended, update to known good version

```
aws eks update-nodegroup-config --cluster-name my-cluster --nodegroup-name my-ng --launch-template version="32" # Previous known good
```

# 3. Wait for recovery

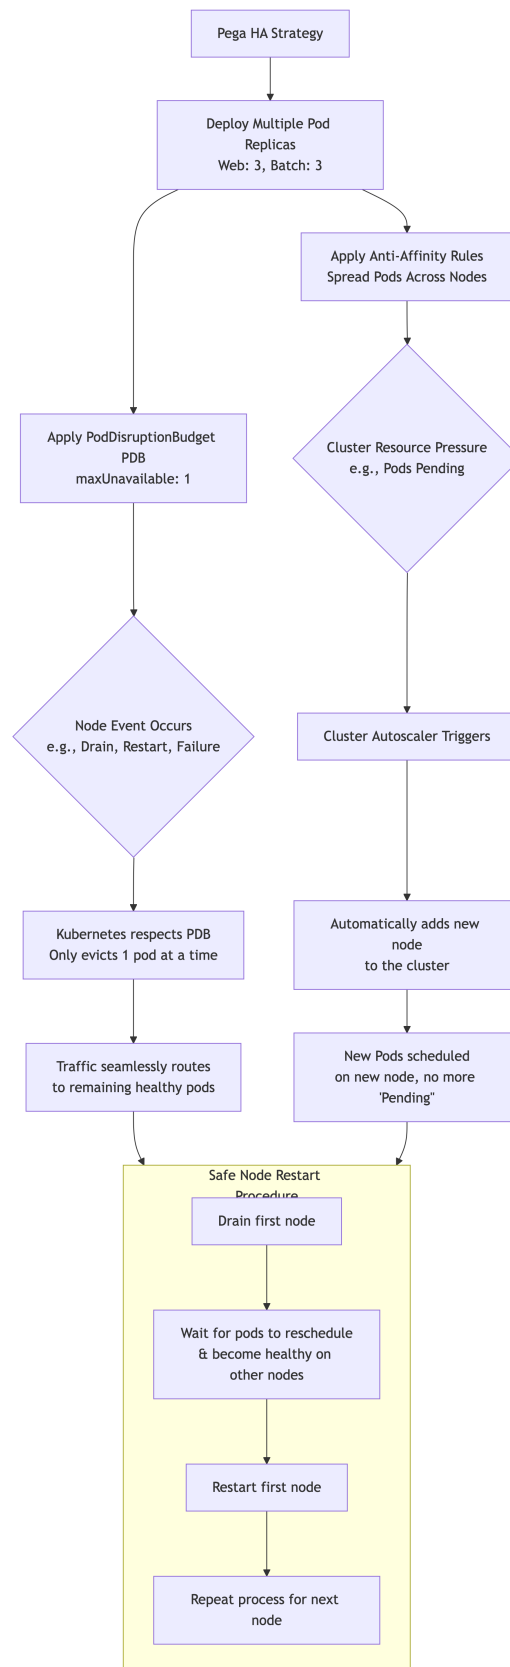
## Workflow of Nodes restart



**EKS HA** : Workflow diagram that visually explains the high-availability strategy, followed by a detailed breakdown of each component.

## High-Level HA Workflow Diagram

This diagram shows the logical flow of how the different Kubernetes concepts work together to achieve high availability

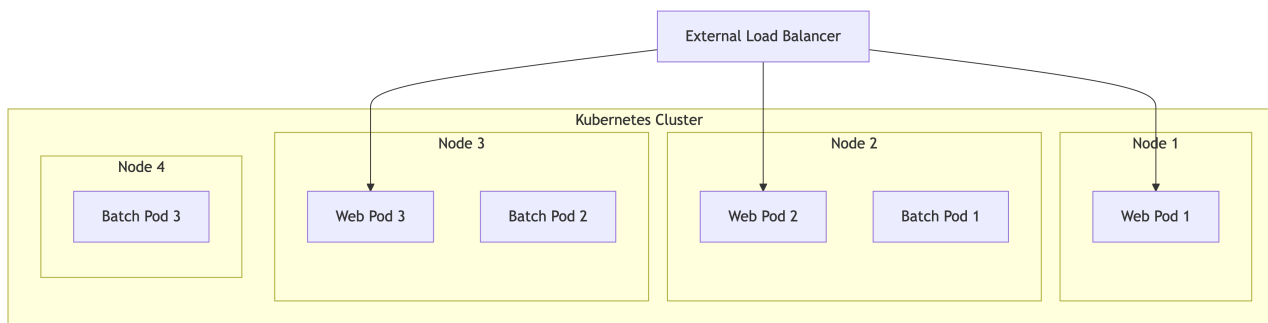


## Detailed Breakdown of Components

The diagram above is built from the following key components working in concert.

### 1. Pod Replicas & Anti-Affinity: The Foundation


This is the base deployment, ensuring multiple copies of your service are running and placed on different physical nodes.



**Key Insight:** By using anti-affinity rules, you ensure no two web pods or two batch pods are on the same node. If **Node 2** fails, you only lose Web Pod 2 and Batch Pod 1. The other web and batch pods on Nodes 1, 3, and 4 continue serving traffic, preventing a full outage.

### 2. PodDisruptionBudget (PDB): The Safety Guard


The PDB acts as a buffer during voluntary disruptions (like node restarts or upgrades).



**Key Insight:** The PDB prevents Kubernetes from accidentally taking down too many pods of a single type at once, ensuring minimum capacity is always maintained during operations.

### 3. Cluster Autoscaler: The Capacity Provider

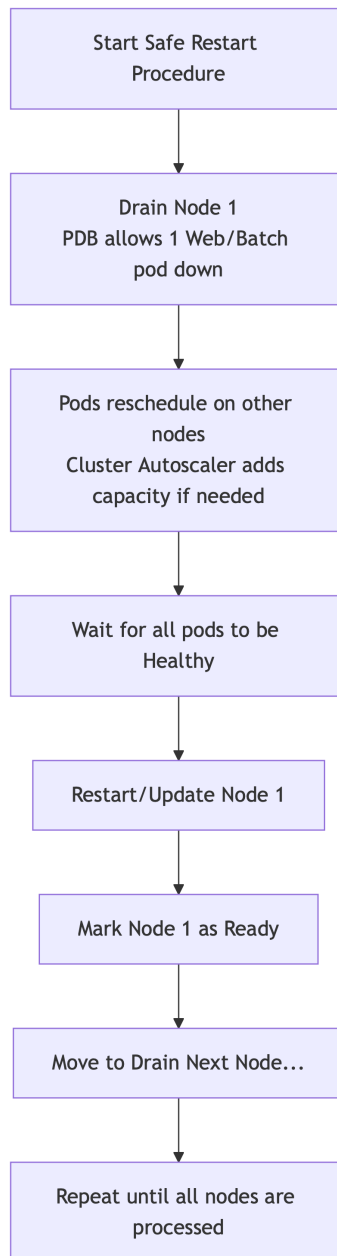
This ensures there is always enough infrastructure for the pods to run on.



**Key Insight:** The autoscaler eliminates the "no room to schedule" failure mode, which is a common cause of downtime in poorly managed clusters.

#### 4. Safe Node Restarts: The Procedure

This is the active process of applying all the above concepts to maintain the cluster without downtime.



**Key Insight:** This deliberate, one-by-one process leverages all the other HA mechanisms (Replicas, PDB, Anti-Affinity, Autoscaler) to guarantee service continuity.

#### One-Line Explanation for the Team

“We keep 3 pods for each service, spread them across nodes, only allow 1 pod down at a time, let the cluster add new nodes when needed, and restart nodes one by one. That’s how we get HA.”



